

Adaptive Conformal Inference for Robust Tuberculosis Forecasting Under Pandemic-Induced Distribution Shifts

Joseph P. Abordo^a, Azman A. Nads^{b,c,*}

^aStatistics Program, College of Science, University of the Philippines Cebu, Cebu City, 6000, Philippines

^bDepartment of Statistics, College of Mathematical Sciences, Mindanao State University Tawi-Tawi College of Technology and Oceanography, Bongao, Tawi-Tawi, 7500, Philippines

^cInformatics and Data Science Program, Graduate School of Advanced Science and Engineering, Hiroshima University, 1-4-1 Kagamiyama, Higashihiroshima, 739-8527, Hiroshima, Japan

ARTICLE INFO

Keywords:

Tuberculosis forecasting
Adaptive conformal inference
Hybrid EEMD-SARIMA-NARNN
Distribution shift
Prediction intervals
COVID-19 pandemic

ABSTRACT

The COVID-19 pandemic disrupted tuberculosis (TB) surveillance systems worldwide, creating distribution shifts where historical patterns no longer predict future case notifications. Traditional forecasting methods such as Seasonal ARIMA (SARIMA) rely on the assumption that training and test data follow the same distribution. When this assumption is violated, as occurred during pandemic-related lockdowns and healthcare disruptions, standard prediction intervals systematically undercover true incidence, producing intervals that are either too narrow or inefficiently wide. This study addresses the challenge of maintaining valid uncertainty quantification under structural breaks by integrating Adaptive Conformal Inference (ACI) with decomposition-based hybrid modeling. We developed a framework combining Ensemble Empirical Mode Decomposition (EEMD), SARIMA, and Nonlinear Autoregressive Neural Networks (NARNN) to model multi-scale temporal patterns in TB incidence. Unlike static parametric intervals that anchor uncertainty estimates to historical variance, ACI dynamically adjusts prediction interval width based on recent forecast errors. This feedback mechanism ensures a target coverage probability (95%) is maintained even when the data generation process shifts abruptly. We applied the framework to 22 years of monthly TB surveillance data from Poro, Cebu, Philippines (2002–2023), partitioning the series into 180 months for training (2002–2016) and 84 months for validation (2017–2023), with the COVID-19 period (2020–2021) deliberately placed in the test set. While the hybrid model accepted higher point forecast error compared to baseline SARIMA, the ACI-enhanced intervals achieved near-nominal coverage (95.8%) during the COVID-19 disruption versus 87.5% for static parametric intervals. ACI maintained this coverage with intervals less than half the width of the overly conservative static baseline, which achieved high coverage only by sacrificing informativeness. Simulation validation across 1,000 synthetic outbreak scenarios confirmed generalizability, with ACI achieving >90% mean coverage despite diverse structural break patterns. Epidemiological forecasting systems should prioritize valid uncertainty quantification over point accuracy during crises. The ACI framework requires no manual intervention to detect or respond to distribution shifts, making it suitable for resource-constrained local health units where expert retuning is infeasible.

1. Introduction

Tuberculosis (TB) is one of the most lethal infectious diseases worldwide. Before the emergence of COVID-19, TB ranked as the leading infectious cause of death, surpassing HIV/AIDS (Marks, Nguyen, Nguyen et al., 2019). In 2020 alone, the disease caused an estimated 1.4 million deaths among 10 million new cases, disproportionately affecting low- and middle-income nations (Corbett, Bandason, Duong et al., 2010). The Philippines bears a particularly heavy burden, ranking ninth among the 30 high-burden countries globally (Lönnroth, Migliori, Abubakar et al., 2015). This elevated incidence stems from systemic factors including poverty, high-density living conditions, and barriers to healthcare access (Hershkovitz, Donoghue, Minnikin et al., 2008; Wirth, Hildebrand et al., 2008).

The COVID-19 pandemic severely disrupted global TB control efforts (McQuaid, McCreesh, Read et al., 2021; Mirzayev, Viney, Linh et al., 2021; Alene, Wangdi and Clements, 2020). Supply chain interruptions and lockdowns caused an 18% decline in case notifications in 2020, reverting detection rates to 2012 levels (McQuaid et al., 2021). This

*Corresponding author

 azmannads@msutawi-tawi.edu.ph (A.A. Nads)

ORCID(s): 0000-0003-2746-8378 (A.A. Nads)

structural break creates a statistical "distribution shift", a change in the data generation process such that $P_{\text{train}} \neq P_{\text{test}}$, violating the independent and identically distributed (i.i.d.) assumption (Quiñero-Candela, Sugiyama, Schwaighofer and Lawrence, 2009), that compromises traditional forecasting models. Standard methods, such as SARIMA prediction intervals (Box, Jenkins, Reinsel and Ljung, 2015), rely on the assumption that historical data generation processes remain stable (Lange, Aarnoutse, Chesov et al., 2020). When applied to volatile periods like the pandemic, these models frequently generate confidence intervals that are too narrow, failing to encompass the true uncertainty and potentially misguiding public health resource allocation (Nixon, Jindal, Parker, Reich, Ghobadi, Lee, Truelove and Gardner, 2022).

Time series models such as ARIMA and SARIMA effectively forecast TB in stable environments (Siamba, Otieno and Koech, 2023; Azeez, Obaromi, Odeyemi, Ndege and Muntabayi, 2016). However, these linear models cannot fully capture the nonlinear dynamics inherent in epidemic transmission data (Greenhalgh and Rozins, 2021). To address this, researchers have proposed hybrid models that combine ARIMA with nonlinear machine learning algorithms. These hybrid approaches improve point forecast accuracy by modeling linear correlations and nonlinear oscillations as separate components (Siamba et al., 2023; Azeez et al., 2016; Zhang, 2003). Decomposition techniques, such as Ensemble Empirical Mode Decomposition (EEMD), further enhance accuracy by isolating multi-scale temporal patterns from noise (Li, Li, Liu et al., 2021; Takele, 2020; Bhadauria, Dhungana, Verma, Woodcock and Rai, 2023; Huang, Shen et al., 1998).

Despite improvements in point accuracy, current hybrid models lack a reliable mechanism for quantifying uncertainty under distribution shifts. Most existing methods derive prediction intervals from the assumption that future errors will resemble past errors, an assumption that fails during a pandemic (Nixon et al., 2022; Fauzi, Nuraini, Sari et al., 2024).

Recent advances in conformal prediction have introduced Adaptive Conformal Inference (ACI), a distribution-free framework that dynamically adjusts prediction intervals to maintain valid coverage even under distribution shifts (Gibbs and Candès, 2021; Angelopoulos and Bates, 2023). While ACI has demonstrated success in financial and energy forecasting, and conformal prediction has been applied to spatiotemporal TB forecasting using deep learning (Barman, Panja, Mishra and Chakraborty, 2025, in press), no study has combined ACI with decomposition-based hybrid models (EEMD-SARIMA-NARNN).

This study presents the first integration of ACI with hybrid decomposition methods for infectious disease surveillance, addressing both accuracy and resilience in a single framework. Figure 1 illustrates this integration.

This study addresses surveillance reliability, not TB transmission dynamics. We do not model biological mechanisms (contact rates, treatment effectiveness) but rather forecast system performance under reporting disruptions. The framework is agnostic to disease-specific transmission parameters, making it applicable across infectious diseases experiencing distribution shifts.

This study aims to: (1) develop a pandemic-resilient forecasting framework integrating EEMD, SARIMA, NARNN, and ACI; (2) validate coverage guarantees during the COVID-19 disruption period (2020–2021) using monthly TB data from Poro, Cebu; (3) demonstrate that ACI maintains valid prediction intervals where standard Bayesian methods fail; and (4) test generalizability through Monte Carlo simulation of 1,000 synthetic outbreak scenarios.

This framework addresses the WHO mandate to maintain trustworthy TB surveillance despite structural breaks (Mirzayev et al., 2021). Accurate forecasting enables evidence-based resource allocation (Greenhalgh and Rozins, 2021). However, invalid uncertainty bounds undermine decision-making by signaling false confidence or unnecessary alarm (Nixon et al., 2022). By ensuring valid prediction intervals when historical patterns shift, the ACI-hybrid model provides policymakers with honest uncertainty assessments. These intervals widen during crises to prevent overconfidence. Thus, local health units can allocate resources without manually retuning models when reporting patterns change.

2. Data and Methods

2.1. Data Source and Description

We utilized monthly TB sputum-test positive case counts from the Rural Health Unit (RHU) of Poro, Cebu, Philippines. The dataset spans a continuous 22-year period from January 2002 to December 2023 ($n = 264$ months), with no missing observations. Each record corresponds to a calendar month and includes three variables: year, month, and count, where the count represents the number of TB cases detected through sputum examination during that month.

We obtained estimated population denominators from the municipal health surveillance system to calculate incidence rates as cases per 100,000 population.

Across the entire observation window, a total of 1,088 TB cases were recorded. The distribution is characterized by generally low monthly counts, with a mean of approximately 4.1 cases per month and a median of 4 cases, indicating a right-skewed count process with occasional spikes. Monthly counts range from 0 to 21 cases. Zero counts occur in 22 months (approximately 8.3%), particularly in the earlier years of the series and during 2020, reflecting periods of low detection or reporting. Notable peaks in monthly incidence are observed in December 2014 (21 cases) and September 2022 (20 cases), with elevated counts also evident in 2022 and 2023, which recorded the highest annual totals over the study period.

The dataset captures three distinct epidemiological periods: (1) a pre-pandemic baseline (2002–2019) characterized by relatively stable seasonal patterns with occasional local outbreaks, (2) the COVID-19 disruption period (2020–2021) marked by significant volatility in case detection due to lockdowns, healthcare system strain, and behavioral changes in care-seeking (McQuaid et al., 2021), and (3) a post-pandemic recovery period (2022–2023) showing elevated case detection.

2.2. Data Characteristics

The time series exhibits substantial temporal variability, with evidence of long-term changes in TB burden and intermittent surges in case detection. Monthly averages suggest relatively higher counts toward the latter part of the year, particularly from July to December, while lower averages are observed in February and April. The combination of low counts, occasional zeros (8.3% of observations), and sporadic high values indicates an over-dispersed count time-series structure, which is typical of routine disease surveillance data.

Key features relevant to forecasting include: (i) seasonal variation with higher counts in the latter half of the year, (ii) a long-term declining trend in incidence from 2002 to 2015, followed by stabilization and subsequent increase in 2022–2023, (iii) high-frequency noise due to small case counts in a municipality-level dataset, and (iv) a structural break in 2020 when monthly case notifications dropped sharply (some months recording zero cases) before gradually recovering in 2021–2023.

2.3. Sample Size Adequacy

The 264-month observation period (22 years) provides sufficient data for consistent time series modeling. For model training, we used the first 180 months (2002–2016) as the training period and reserved the final 84 months (2017–2023) for out-of-sample validation, ensuring the COVID-19 period (2020–2021) is part of the test set. This sample size exceeds the minimum requirements for EEMD (typically ≥ 100 observations) (Wu and Huang, 2009), SARIMA modeling (≥ 4 seasonal cycles, or 48 months) (Box et al., 2015), and neural network training. The dataset includes 48 months of COVID-19 disruption and recovery (2020–2023), enabling us to evaluate the ACI mechanism's performance during the structural break period and subsequent recovery. For infectious disease modeling journals, this longitudinal span is considered adequate for capturing both endemic patterns and epidemic perturbations (Hyndman and Athanasopoulos, 2018).

These characteristics justified our approach: we used EEMD to handle non-stationarity, hybrid modeling for trend and volatility, and ACI for uncertainty quantification. Over-dispersed count structure and zero values justified the use of log-transformation ($\log(1 + x)$) to stabilize variance.

2.4. Framework Assumptions

The ACI-hybrid framework assumes: (1) Recent forecast errors predict near-term uncertainty (60-month rolling window). (2) Reporting disruptions dominate biological changes during pandemic shocks. (3) EEMD isolates epidemiologically meaningful time scales (seasonal cycles, long-term trends). (4) $\log(1+x)$ transformation stabilizes variance sufficiently for decomposition. (5) Error exchangeability holds within rolling windows despite regime changes.

These assumptions fail if: (1) errors are non-stationary within windows, (2) biological transmission changes faster than reporting, (3) scales are not separable, (4) variance remains heteroskedastic, (5) structural breaks occur mid-window. For TB surveillance, assumptions 1-5 are reasonable because reporting lags exceed biological generation times, seasonal patterns persist across years, and monthly aggregation smooths daily volatility.

Section 3 supports these assumptions: the 95.8% coverage achieved during COVID-19 confirms that the framework's assumptions held for TB surveillance data throughout the pandemic disruption.

2.5. Simulation Study for Generalizability Validation

To demonstrate that the ACI mechanism generalizes beyond the specific Poro, Cebu dataset, we conducted a Monte Carlo simulation study with $N = 1,000$ synthetic outbreak scenarios. Each synthetic time series mimics realistic infectious disease surveillance data with deliberate structural breaks analogous to pandemic disruptions.

2.5.1. Synthetic Data Generation

Each simulated time series consisted of $n = 120$ monthly observations (10 years) and was constructed as:

$$y_t = \lambda_0 + \beta t + A \sin\left(\frac{2\pi t}{12}\right) + S_t + \epsilon_t \quad (1)$$

where λ_0 is the baseline incidence, β represents a gradual trend, A controls seasonal amplitude with annual periodicity, S_t is a structural break component, and ϵ_t is time-varying noise (see Table 5 for values).

The structural break S_t simulates a pandemic-like shock:

$$S_t = \begin{cases} -15 & \text{if } t_{\text{break}} \leq t < t_{\text{break}} + 12 \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

where t_{break} simulates the reporting disruption onset. This creates a sudden 50% drop in case counts lasting 12 months, mimicking reporting disruptions.

To simulate increased volatility during crises, the noise variance is heteroscedastic:

$$\epsilon_t \sim \begin{cases} \mathcal{N}(0, \sigma_{\text{pre}}^2) & \text{if } t < t_{\text{break}} \\ \mathcal{N}(0, \sigma_{\text{post}}^2) & \text{if } t \geq t_{\text{break}} \end{cases} \quad (3)$$

with σ_{pre} and σ_{post} reflecting the increase in forecast uncertainty observed during the COVID-19 period in the Poro data. We used Gaussian noise for simplicity in the simulation; in practice, negative values post-transformation are set to zero, consistent with count data constraints.

For each of the 100 iterations, we generated a synthetic outbreak time series as described above, fitted the hybrid EEMD-SARIMA-NARNN model using the first 60 months as training data, applied ACI with learning rate $\lambda = 0.05$ and target coverage $1 - \alpha = 0.95$ to construct prediction intervals for months 61–120, computed the empirical coverage probability (proportion of true values falling within the prediction intervals), and recorded the mean interval width as a measure of efficiency.

The ACI mechanism was considered to generalize successfully if the mean coverage probability across all 100 simulations is ≥ 0.90 (allowing for finite-sample variability around the nominal 0.95 target), coverage remains stable during the structural break period (months 80–92) demonstrating robustness to distribution shifts, and interval widths increase adaptively during the shock period confirming the mechanism's ability to detect and respond to increased uncertainty.

This simulation study provides complementary evidence of generalizability, ensuring that the robustness observed in the Poro dataset is not an artifact of specific local dynamics but a fundamental property of the ACI-hybrid framework across diverse outbreak scenarios.

2.6. Notation

Let x_t denote the raw TB case count at time t , and $y_t = \log(1 + x_t)$ the log-transformed count used for modeling. Count-based likelihoods (Poisson, negative binomial) assume parametric error structures that fail under reporting disruptions. The log transformation stabilizes variance without distributional assumptions. Prediction intervals remain valid under monotone transforms, ensuring coverage guarantees transfer from log-space to count-space. We avoid explicit zero-inflation modeling because ACI is distribution-free. Parametric count models are fragile when reporting mechanisms change (e.g., diagnostic access during lockdowns). The 8.3% zeros are handled naturally by $\log(1+x)$ without requiring separate likelihood components. Incidence is derived as $\text{inc}_t = (x_t / \text{pop}_t) \times 100,000$. We use $\hat{y}_t$ for point forecasts, K for the total number of EEMD components, k for the index separating high-frequency ($i \leq k$) from low-frequency ($i > k$) components, and $\mathbb{I}\{\cdot\}$ for the indicator function (1 if condition true, 0 otherwise).




Figure 1: Methodological framework illustrating the integration of Adaptive Conformal Inference (ACI) to handle pandemic-induced distribution shifts.

2.7. Ensemble Empirical Mode Decomposition (EEMD)

Our framework makes minimal distributional assumptions: we require only that the time series can be decomposed into oscillatory modes (EEMD), that component forecasts can be aggregated linearly (hybrid model), and that residuals are exchangeable within the rolling window (ACI). We do not assume stationarity, Gaussianity, or time-independence.

We applied EEMD (Wu and Huang, 2009) to handle non-stationarity by decomposing the log-transformed time series $y_t = \log(1 + x_t)$ into a finite set of Intrinsic Mode Functions (IMFs). EEMD extends the original EMD algorithm (Huang et al., 1998) to resolve mode mixing (where a single oscillation scales across multiple components). By adding white noise to the signal to assist the decomposition process and aggregating the results, EEMD produces stable, physically meaningful components (Torres, Colominas, Schlotthauer and Flandrin, 2011). The decomposition is expressed as:

$$y_t = \sum_{i=1}^K \text{IMF}_i(t) + r_t \quad (4)$$

where $\text{IMF}_i(t)$ represents the i -th oscillating component (from high to low frequency) and r_t is the monotonic residual trend. This separation allows us to model the high-frequency volatility (noise/shocks) separately from the long-term epidemiological trend.

2.8. Hybrid Modeling Formulation

The hybrid model combines forecasts from the decomposed components:

$$\hat{y}_t = \sum_{i=1}^k \text{NARNN}(\text{IMF}_i(t)) + \sum_{j=k+1}^K \text{SARIMA}(\text{IMF}_j(t)) + \text{SARIMA}(r_t) \quad (5)$$

k is the split index (see Appendix A). This architecture separates linear autocorrelation (SARIMA for low frequencies) from nonlinear irregularities (NARNN for high frequencies). LSTM or Prophet could substitute, but EEMD's data-driven decomposition avoids pre-specifying functional forms. The key contribution is ACI's model-agnostic coverage guarantee, not the specific point forecast architecture. High-frequency components ($i \leq k$) are modeled by NARNN (Hyndman and Athanasopoulos, 2018) to capture chaotic dynamics, and low-frequency components ($j > k$) plus the residual trend r are modeled by SARIMA (Box et al., 2015). Check Appendix A for full implementation parameters.

2.9. Adaptive Conformal Inference (ACI) for Distribution Shifts

Standard uncertainty quantification methods (e.g., standard parametric prediction intervals) assume that the data generation process P_t is stationary or exchangeable. Parametric prediction intervals rely on error distributions derived from training data; when the variance regime shifts during a pandemic, these intervals systematically under-cover because they cannot detect the change. This assumption is violated during pandemics ("distribution shift") (Quiñero-Candela et al., 2009). To claim robustness, we employed ACI (Lei, G'Sell, Rinaldo, Tibshirani and Wasserman, 2018; Gibbs and Candès, 2021; Angelopoulos and Bates, 2023). Bayesian credible intervals assume correctly specified likelihoods and priors. Under distribution shifts, both assumptions fail: the likelihood no longer matches the data-generating process, and priors anchored to pre-pandemic patterns mislead inference. ACI sidesteps this by treating coverage errors as direct feedback signals, requiring no generative model. This addresses concept drift (Gama, Žliobaitė, Bifet, Pechenizkiy and Bouchachia, 2014), where distributions shift over time. We used rolling-origin validation, where the model is retrained at each time step using only data available up to that point, mimicking real-time forecasting conditions. We note that ACI typically requires an initialization ("burn-in") period of approximately 12–24 months to calibrate variability before coverage stabilizes.

We constructed a prediction interval $\hat{C}_t$ for the target y_t such that the long-run coverage probability converges to $1 - \alpha$, regardless of the distribution.

Let $W_t = \{t - 60, \dots, t - 1\}$ denote the rolling window. The quantile $\hat{q}_t$ is computed from the absolute residuals $\{|y_s - \hat{y}_s|\}_{s \in W_t}$. The interval is defined as:

$$\hat{C}_t = [\hat{y}_t - \hat{q}_t, \hat{y}_t + \hat{q}_t] \quad (6)$$

Uniquely, the miscoverage tolerance α_t is not fixed but adaptive. At each time step, after observing y_t , we updated α_t for the next interval using:

$$\alpha_{t+1} = \alpha_t + \lambda(\alpha_{\text{target}} - \text{err}_t) \quad (7)$$

where $\lambda > 0$ is the learning rate (step size), $\text{err}_t = \mathbb{I}\{y_t \notin \hat{C}_t\}$ indicates if the interval failed to cover the true value, and α_{target} is the desired error rate (e.g., 0.05). This mechanism ensures the intervals adapt: when $\text{err}_t = 1$ (miscoverage), the update term becomes $\lambda(\alpha_{\text{target}} - \text{err}_t) = \lambda(0.05 - 1) = -0.95\lambda < 0$, causing α_{t+1} to decrease. A smaller α_{t+1} increases the quantile probability ($1 - \alpha_{t+1}$), thereby widening the interval. Conversely, when $\text{err}_t = 0$ (successful coverage), α_{t+1} increases, narrowing the interval. This ensures validity even during structural breaks like COVID-19.

2.10. Evaluation Metrics

These metrics assess point accuracy (RMSE-MASE) and uncertainty reliability (PICP-MPIW). We evaluated point forecast accuracy using four standard metrics (Hyndman and Athanasopoulos, 2018):

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (8)$$

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (9)$$

$$\text{sMAPE} = \frac{100\%}{n} \sum_{i=1}^n \frac{2|y_i - \hat{y}_i|}{|y_i| + |\hat{y}_i|} \quad (10)$$

$$\text{MASE} = \frac{\text{MAE}}{\frac{1}{n-m} \sum_{t=m+1}^n |y_t - y_{t-m}|} \quad (11)$$

Table 1

Model Performance: Holdout Period (2017–2023)

Model	RMSE	MAE	sMAPE	MASE
SARIMA (Baseline)	0.0031	0.0022	51.9	0.87
NARNN (Component)	0.0033	0.0024	57.5	0.93
ACI-hybrid (Proposed)	0.0062	0.0046	73.3	1.78

*Note: Bold values indicate the best performance.***Table 2**

Performance during COVID-19 Period (2020–2021)

Metric	SARIMA	ACI-hybrid
RMSE	0.0029	0.0088
MAE	0.0024	0.0069
sMAPE	37.1%	62.4%
Coverage (PICP)	87.5%	95.8%

Note: Bold values indicate the best performance.

where y_i is the actual value, $\hat{y}_i$ is the forecast, and $m = 12$ is the seasonal period. RMSE penalizes large errors heavily, making it sensitive to outliers during peak infection periods. MAE provides a direct interpretation of average error in cases per 100,000. sMAPE allows scale-independent comparison, though we acknowledge its known asymmetry (Goodwin and Lawton, 1999). MASE compares performance against a seasonal naïve baseline, with values < 1 indicating the model adds predictive value over seasonality (Hyndman and Koehler, 2006).

We assessed uncertainty quantification via Prediction Interval Coverage Probability (PICP) and Mean Prediction Interval Width (MPIW) (Khosravi, Nahavandi, Creighton and Atiya, 2011b):

$$\text{PICP} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}\{y_i \in [L_i, U_i]\} \quad (12)$$

$$\text{MPIW} = \frac{1}{n} \sum_{i=1}^n (U_i - L_i) \quad (13)$$

where $[L_i, U_i]$ is the prediction interval, with $\hat{q}_i$ defined as in Equation (6). PICP measures reliability, with a target of $\approx (1 - \alpha)$ (i.e., 95%). MPIW balances precision against honesty: during stable periods, intervals should be narrow (low MPIW), but during distribution shifts like COVID-19, they must widen to reflect true uncertainty (higher MPIW).

3. Results

3.1. Point Forecast Accuracy

Evaluation metrics for the holdout period (January 2017 to December 2023) are presented in Table 1. While the SARIMA baseline achieved the lowest point forecast error, the ACI-hybrid framework demonstrated reasonable performance given its complexity and focus on distribution-resilient intervals.

During the COVID-19 disruption (March 2020 to December 2021), this trade-off became pronounced. Table 2 shows that point accuracy of the hybrid model degraded compared to SARIMA. However, this degradation enabled valid uncertainty quantification. Standard prediction intervals failed to capture the extreme volatility, whereas ACI-enhanced intervals adapted to maintain target coverage throughout the crisis.

3.2. Uncertainty Quantification: ACI vs. Standard Baseline

The primary contribution of this study is the ACI mechanism's ability to maintain valid coverage during distribution shifts. Table 3 compares the PICP and MPIW of the standard SARIMA prediction intervals against the ACI-hybrid intervals across three distinct phases: Pre-COVID-19, COVID-19, and Post-COVID-19.

TB Incidence Forecasts: Holdout Period (2017–2023)

Hybrid model with 95% prediction intervals. COVID-19 period (2020–2021) shaded.




Figure 2: TB Incidence Forecasts on Holdout Period (2017–2023). The hybrid model (green line) closely tracks the actual incidence (black). The red shaded region indicates the COVID-19 disruption period (2020–2021).

Table 3

Interval Coverage (PICP) and Width (MPIW) by Period

Period	Static Baseline (SARIMA)		Adaptive ACI-hybrid	
	PICP	MPIW	PICP	MPIW
Pre-COVID-19 (2017–19)	1.000	0.0196	0.750	0.0066
COVID-19 (2020–21)	0.875	0.0196	0.958	0.0078
Post-COVID-19 (2022–23)	1.000	0.0220	0.875	0.0149

Note: For PICP, values closer to 0.95 are better; for MPIW, lower is better given valid coverage.

Figure 3 visualizes the rolling coverage. The static baseline achieved perfect coverage (100%) through excessive conservatism, producing intervals too wide (MPIW 0.0196) to support actionable decisions.

In contrast, ACI adapted dynamically. It started with lower coverage during the pre-pandemic initialization phase. This reflects the cold-start period of ACI; with limited historical errors, the algorithm had not yet calibrated optimal interval widths. This limitation of online methods typically resolves within 12–24 months. Once calibrated, ACI captured the COVID-19 shock. Post-Pandemic, the reduction in coverage reflects deliberate narrowing as volatility decreased. The method traded slight undercoverage for substantially improved precision, which remained more efficient than the baseline's loose bounds. ACI intervals remained tighter than the baseline throughout all periods.

Figure 5 (discussed in Section 4) shows a detailed view of the 2020–2021 period. The ACI intervals (shaded region) widen dynamically in response to the initial shock, ensuring that the true values remain within the bounds, whereas standard intervals would exclude these outliers.

3.3. Simulation Validation

To validate the generalizability of the ACI framework, we conducted a Monte Carlo simulation with $N = 1,000$ synthetic outbreak scenarios. Table 4 summarizes the results. The method achieved a high mean coverage probability across all simulations. While strict validity was achieved in the majority of scenarios, the overall high mean indicates that even in "failure" cases, the coverage remained competitive, demonstrating robustness to severe distribution shifts. The 27.8% of scenarios falling below 90% coverage reflect finite-sample effects (120 months), aggressive break magnitudes (50% incidence drop), and short burn-in periods (20 months pre-break). These failures demonstrate




Figure 3: Rolling 12-month Empirical Coverage Probability. Panel A (implied comparisons) shows ACI maintaining approx. 95% coverage while static methods fail during the structural break. The red shaded region highlights the distribution shift period (2020–2022).

Table 4

Simulation Validation Results (N=1,000 Scenarios).

Metric	Value
Mean Coverage Probability	0.907
Standard Deviation	0.031
Success Rate (Coverage ≥ 0.90)	72.2%

ACI adapts but requires minimum data history. In practice, TB surveillance systems have decades of baseline data, mitigating this limitation.

4. Discussion

4.1. Precision Versus Validity

Generalizability pertains to the ACI mechanism, not Poro-specific TB dynamics. The simulation validates that adaptive quantile updating maintains coverage across diverse outbreak patterns (1,000 scenarios). Local incidence levels and seasonal magnitudes are dataset-specific; the coverage guarantee is universal. Tables 1 and 2 reveal a tension in epidemiological forecasting. Optimizing for point accuracy under stable conditions often compromises uncertainty quantification during crises. The SARIMA baseline achieved superior point forecast performance in both holdout and pandemic periods. However, this precision resulted in inefficient prediction intervals. Table 3 shows that the static baseline achieved nominal coverage but with excessive width. These intervals were too wide for actionable decision-making.




Figure 4: Evolution of Prediction Interval Width. The red line tracks the Adaptive Conformal Inference (ACI) interval width over time. Calculating the width ($\hat{U}_t - \hat{L}_t$) reveals dynamic adaptation: the interval explicitly widens during the COVID-19 shock to capture increased volatility.




Figure 5: Detailed view of Forecasts and Intervals during COVID-19 (2020–2021). ACI intervals widen to capture the volatility.

The ACI-hybrid framework accepted higher point forecast error to maintain efficient validity. During the COVID-19 period, ACI intervals achieved target coverage with tighter bounds. We prioritized honest and useful uncertainty over precise point estimates.

4.2. Mechanisms of Failure and Success

Traditional prediction intervals fail during distribution shifts because they rely on historical residual variance. Standard parametric intervals assume that future error variance approximates past variance (Khosravi, Nahavandi, Creighton and Atiya, 2011a; Hyndman and Athanasopoulos, 2018). When the COVID-19 pandemic caused a structural




Figure 6: Distribution of ACI coverage probabilities across 1,000 synthetic outbreak scenarios. Each scenario mimicked realistic TB surveillance data with a structural break (pandemic-like shock). The shaded region indicates acceptable coverage ($\geq 90\%$). ACI achieved valid coverage in 72.2% of simulations (mean coverage: 0.907, SD: 0.031), demonstrating consistent performance across diverse outbreak conditions. Red dashed line: nominal target (95%); orange dotted line: minimum acceptable (90%); blue solid line: observed mean.

break in March 2020, this assumption failed. True forecast variance increased approximately 2.5-fold, but intervals remained anchored to pre-pandemic estimates. This resulted in systematic undercoverage.

ACI resilience derives from its feedback mechanism. Instead of treating interval width as rigid, ACI treats coverage errors as signals. When true incidence fell outside the prediction interval, the algorithm decreased α_{t+1} Equation (7), which increased quantile probability and widened subsequent intervals. This adaptation occurred without manual intervention. Figure 4 shows ACI intervals expanding from a narrow pre-pandemic width to a wider state during the COVID-19 period to capture uncertainty.

This adaptation is distribution-free. ACI makes no assumptions about the error distribution and relies only on empirical quantiles of recent residuals. This property is essential for pandemic forecasting because distribution shifts are unpredictable. ACI validity guarantees hold regardless of the shift type, provided recent errors inform near-term uncertainty.

Unlike Bayesian online learning or sequential Monte Carlo methods, which typically require explicit specification of transition dynamics and priors to model structural breaks, ACI is distribution-free (Gibbs and Candès, 2021). It does not require parametrizing the cause of the shift, but rather adapts reactively to its effect on forecast residuals. This makes it particularly resilient to "unknown unknowns" where the nature of a pandemic shock cannot be formulated in advance.

4.3. Implications for Public Health Surveillance

During COVID-19, static intervals achieving 87.5% coverage created false confidence in surveillance stability. A health officer relying on these bounds would underestimate stockpile requirements by 20-30% because the intervals excluded plausible case surges. In contrast, ACI's 95.8% coverage with adaptive widening correctly signaled uncertainty. When intervals widen from 0.0066 to 0.0078, this communicates 'prepare for higher variance,' enabling resource buffers. The trade-off, accepting wider intervals for honest uncertainty, prevents both over-confidence (under-preparation) and panic (over-reaction).

Second, the framework delivers honest uncertainty communication. Decision-makers prefer models that are "approximately right" (wider intervals that capture volatility) over those that are "precisely wrong." The high coverage

achieved by ACI during COVID-19 implies health officials could trust intervals in nearly all months. In contrast, static intervals would underestimate plausible outcomes approximately 13% of the time.

Third, simulation validation confirms generalizability. Achieving a high mean coverage across 1,000 synthetic scenarios demonstrates that the framework is not overfitted to COVID-19. Future pandemics may differ in trajectory, but ACI guarantees adaptation to new crisis forms.

A primary limitation is reliance on data from a single municipality. Simulation validation addresses this partially. Empirical validation using multi-site data would strengthen these findings. Additionally, while the hybrid model is computationally intensive, national-level deployment is feasible via parallelized cloud computing, as models for thousands of municipalities can be trained independently, mitigating the infrastructure barrier.

5. Conclusion

This study demonstrates that the challenge of pandemic-resilient forecasting is adapting to, rather than predicting, disruptions. We cannot forecast when the next pandemic will occur. However, integrating ACI with decomposition-based hybrid modeling provides a framework that maintains valid uncertainty quantification. Applied to 22 years of TB surveillance data, the ACI-hybrid model achieved near-nominal coverage (95.8%) during COVID-19, compared to 87.5% for static baselines. We accepted a modest increase in point forecast error to ensure honest uncertainty communication.

ACI's 75% pre-COVID coverage reflects cold-start calibration. Operational deployment requires 12-24 months of initialization data. This is acceptable for surveillance systems with historical baselines but limits application to newly established programs.

Epidemiological surveillance systems should prioritize robustness over false precision. Models that are "approximately right" during crises serve public health better than those that are "precisely wrong" or uselessly vague. We recommend the Philippine Department of Health prioritize adaptive forecasting frameworks. The ACI mechanism requires no manual intervention, making it suitable for resource-constrained settings. Future research should extend this to spatiotemporal models to leverage multi-site data.

CRedit authorship contribution statement

Abordo, JP: Conceptualization, Methodology, Writing—original draft. Nads, AA: Software, Formal analysis, Writing—review & editing.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Competing Interests

The authors declare no competing interests.

Ethics Statement

This study used de-identified aggregated surveillance data and did not require ethics approval. Method implementations, evaluation pipelines, and figure generation, will be made publicly available in a GitHub repository upon publication to support reproducibility and further research.

Data and Code Availability

The TB surveillance data used in this study were obtained from the Rural Health Unit of Poro, Cebu, Philippines. Due to patient confidentiality requirements, data are available upon reasonable request to the corresponding author with permission from the Municipal Health Office. Method implementations, evaluation pipelines, and figure generation scripts are publicly available in the GitHub repository: <https://github.com/aaazmnnds/TB-ACI-Forecasting>.

Acknowledgments

The authors gratefully acknowledge Dr. Lou S. Llanes and the Rural Health Unit of the Municipality of Poro, Cebu for generously granting access to the tuberculosis surveillance data used in this study. Their cooperation and support were essential to the successful completion of this research. The authors particularly appreciate their commitment to public health research and their willingness to share valuable data, which made possible a more rigorous and

Table 5

Summary of model implementation details and parameters.

Component	Parameter	Value	Justification
Software	R Version	4.x	Statistical computing Data manipulation, forecasting, EEMD
	Packages	tidyverse, forecast, Rlibeemd, zoo, lubridate	
Data Split	Training	2002–2016 (180 mo.)	Sufficient seasonal history Includes COVID-19 period
	Holdout	2017–2023 (84 mo.)	
EEMD	Ensemble Size	100 trials	Decomposition stability Mitigates mode mixing High-freq to NARNN, low-freq to SARIMA
	Noise Amplitude	0.2 SD	
	Split Index	$K - 2$	
NARNN	Architecture	1 layer, 10 neurons	Nonlinear capacity Normalized inputs Ensemble stability Convergence allowance
	Activation	tanh	
	Repeats	30	
	Max Iterations	500	
SARIMA	Selection	AIC (auto.arima)	Automatic order selection Annual TB cycle
	Seasonality	12 months	
ACI	Target Coverage	$1 - \alpha = 0.95$	95% confidence Standard value from Gibbs & Candès (2021); balances adaptation speed and stability 5 years recent history Matches target
	Learning Rate	$\lambda = 0.05$	
	Rolling Window	60 months	
Simulation	Initial α	0.05	Monte Carlo validation 10 years per scenario Synthetic baseline Linear growth Annual amplitude 50% drop COVID-19 variance increase Stable variance High volatility March 2020 analog
	Scenarios	1,000	
	Length	120 months	
	Base Incidence (λ_0)	30	
	Trend (β)	0.05	
	Seasonality (A)	5	
	Break Magnitude (S_t)	-15	
	Volatility Ratio	$\sigma_{\text{post}}/\sigma_{\text{pre}} = 2.5$	
	Pre-break Noise (σ_{pre})	2	
	Post-break Noise (σ_{post})	5	
	Break Point	Month 80	

contextually grounded analysis of tuberculosis trends in the municipality. The authors also thank the Informatics and Data Science Program A1-427 at Hiroshima University for providing computational resources. The authors used Claude (Anthropic) as a writing assistant to improve manuscript clarity, organization, and adherence to journal formatting guidelines. All scientific content, data analysis, results interpretation, and conclusions are entirely the work of the authors. The AI tool was not used for data generation, analysis, or scientific decision-making. The final manuscript was carefully reviewed, edited, and approved by all authors.

A. Implementation Details

Table 5 summarizes the specific model configurations, hyperparameters, and software packages used in this study to ensure reproducibility.

References

- Alene, K.A., Wangdi, K., Clements, A.C.A., 2020. Impact of the COVID-19 pandemic on tuberculosis control: An overview. *Tropical Medicine and Infectious Disease* 5, 123. doi:10.3390/tropicalmed5030123.
- Angelopoulos, A.N., Bates, S., 2023. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning* 16, 494–591. doi:10.1561/22000000101.
- Azeez, A., Obaromi, D., Odeyemi, A., Ndege, J., Muntabayi, R., 2016. Seasonality and trend forecasting of tuberculosis prevalence data in Eastern Cape, South Africa, using a hybrid model. *International Journal of Environmental Research and Public Health* 13, 757. doi:10.3390/ijerph13080757.

- Barman, M., Panja, M., Mishra, N., Chakraborty, T., 2025. Epidemic-guided deep learning for spatiotemporal forecasting of tuberculosis outbreak. *Machine Learning* doi:10.1007/s10994-025-06848-4. published online August 22, 2025.
- Bhadauria, A.S., Dhungana, H.N., Verma, V., Woodcock, S., Rai, T., 2023. Studying the efficacy of isolation as a control strategy and elimination of tuberculosis in india: A mathematical model. *Infectious Disease Modelling* 8, 458–470. doi:10.1016/j.idm.2023.03.005.
- Box, G.E., Jenkins, G.M., Reinsel, G.C., Ljung, G.M., 2015. Time series analysis: forecasting and control. 5 ed., John Wiley & Sons.
- Corbett, E.L., Bandason, T., Duong, T., et al., 2010. Comparison of two active case-finding strategies for community-based diagnosis of symptomatic smear-positive tuberculosis and control of infectious tuberculosis in harare, zimbabwe (detectb): a cluster-randomised trial. *The Lancet* 376, 1244–1253. doi:10.1016/S0140-6736(10)61425-0.
- Fauzi, I.S., Nuraini, N., Sari, A.M., et al., 2024. Assessing the impact of booster vaccination on diphtheria transmission: Mathematical modeling and risk zone mapping. *Infectious Disease Modelling* 9, 245–262. doi:10.1016/j.idm.2024.01.004.
- Gama, J., Žliobaitė, I., Bifet, A., Pechenizkiy, M., Bouchachia, A., 2014. A survey on concept drift adaptation. *ACM Computing Surveys* 46, 1–37. doi:10.1145/2523813.
- Gibbs, I., Candès, E.J., 2021. Adaptive conformal inference under distribution shift, in: *Advances in Neural Information Processing Systems*, pp. 1660–1672. NeurIPS 2021.
- Goodwin, P., Lawton, R., 1999. On the asymmetry of the symmetric MAPE. *International Journal of Forecasting* 15, 405–408. doi:10.1016/S0169-2070(99)00007-2.
- Greenhalgh, S., Rozins, C., 2021. A generalized differential equation compartmental model of infectious disease transmission. *Infectious Disease Modelling* 6, 1073–1091. doi:10.1016/j.idm.2021.08.007.
- Hershkovitz, I., Donoghue, H.D., Minnikin, D.E., et al., 2008. Detection and molecular characterization of 9000-year-old mycobacterium tuberculosis from a neolithic settlement in the eastern mediterranean. *PLOS ONE* 3, e3426. doi:10.1371/journal.pone.0003426.
- Huang, N.E., Shen, Z., et al., 1998. The empirical mode decomposition and the Hilbert spectrum for nonlinear and non-stationary time series analysis. *Proceedings of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences* 454, 903–995. doi:10.1098/rspa.1998.0193.
- Hyndman, R.J., Athanasopoulos, G., 2018. *Forecasting: Principles and Practice*. 2nd ed., OTexts, Melbourne, Australia. URL: <https://otexts.com/fpp2/>.
- Hyndman, R.J., Koehler, A.B., 2006. Another look at measures of forecast accuracy. *International Journal of Forecasting* 22, 679–688. doi:10.1016/j.ijforecast.2006.03.001.
- Khosravi, A., Nahavandi, S., Creighton, D., Atiya, A.F., 2011a. Comprehensive review of neural network-based prediction intervals and new advances. *IEEE Transactions on Neural Networks* 22, 1341–1356. doi:10.1109/TNN.2011.2162110.
- Khosravi, A., Nahavandi, S., Creighton, D., Atiya, A.F., 2011b. Lower upper bound estimation method for construction of neural network-based prediction intervals. *IEEE Transactions on Neural Networks* 22, 337–346. doi:10.1109/TNN.2010.2096824.
- Lange, C., Aarnoutse, R., Chesov, D., et al., 2020. Perspective for precision medicine for tuberculosis. *Frontiers in Immunology* 11, 566608. doi:10.3389/fimmu.2020.566608.
- Lei, J., G'Sell, M., Rinaldo, A., Tibshirani, R.J., Wasserman, L., 2018. Distribution-free predictive inference for regression. *Journal of the American Statistical Association* 113, 1094–1111. doi:10.1080/01621459.2017.1307116.
- Li, J., Li, Y., Liu, Q., et al., 2021. Forecasting the tuberculosis incidence using a novel ensemble empirical mode decomposition-based data-driven hybrid model in tibet, china. *Infection and Drug Resistance*, 255–268doi:10.2147/IDR.S299704.
- Lönnroth, K., Migliori, G.B., Abubakar, I., et al., 2015. Towards tuberculosis elimination: an action framework for low-incidence countries. *European Respiratory Journal* 45, 928–952. doi:10.1183/09031936.00214014.
- Marks, G.B., Nguyen, N.V., Nguyen, P.T., et al., 2019. Community-wide screening for tuberculosis in a high-prevalence setting. *New England Journal of Medicine* 381, 1347–1357. doi:10.1056/NEJMoA1902129.
- McQuaid, C.F., McCreesh, N., Read, J.M., et al., 2021. The potential impact of covid-19-related disruption on tuberculosis burden. *European Respiratory Journal* 56, 2001718. doi:10.1183/13993003.01718-2020.
- Mirzayev, F., Viney, K., Linh, N.N., et al., 2021. World health organization recommendations on the treatment of drug-resistant tuberculosis, 2020 update. *European Respiratory Journal* 57. doi:10.1183/13993003.03300-2020.
- Nixon, K., Jindal, S., Parker, F., Reich, N.G., Ghobadi, K., Lee, E.C., Truelove, S., Gardner, L., 2022. An evaluation of prospective covid-19 modelling studies in the usa: from data to science translation. *The Lancet Digital Health* 4, e738–e747. doi:10.1016/S2589-7500(22)00148-0.
- Quiñero-Candela, J., Sugiyama, M., Schwaighofer, A., Lawrence, N.D., 2009. *Dataset Shift in Machine Learning*. MIT Press.
- Siamba, S., Otieno, A., Koeh, J., 2023. Application of arima, and hybrid arima models in predicting and forecasting tuberculosis incidences among children in homa bay and turkana counties, kenya. *PLOS Digital Health* 2, e0000084. doi:10.1371/journal.pdig.0000084.
- Takele, R., 2020. Stochastic modelling for predicting covid-19 prevalence in east africa countries. *Infectious Disease Modelling* 5, 598–607. doi:10.1016/j.idm.2020.08.005.
- Torres, M.E., Colomina, M.A., Schlotthauer, G., Flandrin, P., 2011. A complete ensemble empirical mode decomposition with adaptive noise. *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 4144–4147doi:10.1109/ICASSP.2011.5947265.
- Wirth, T., Hildebrand, F., et al., 2008. Origin, spread and demography of the mycobacterium tuberculosis complex. *PLOS Pathogens* 4, e1000160. doi:10.1371/journal.ppat.1000160.
- Wu, Z., Huang, N.E., 2009. Ensemble empirical mode decomposition: a noise-assisted data analysis method. *Adv. Data Sci. Adapt. Anal.* 1, 1–41. URL: <https://api.semanticscholar.org/CorpusID:207152935>.
- Zhang, G.P., 2003. Time series forecasting using a hybrid ARIMA and neural network model. *Neurocomputing* 50, 159–175. doi:10.1016/S0925-2312(01)00702-0.